

Flight Price Prediction with XGBoost

August 25, 2024

1 Business Understanding

Dataset contains information about flight booking options from the website Easemytrip for flight travel between India's top 6 metro cities. There are 300261 datapoints and 11 features in the cleaned dataset. Our goal is to (1) answer the business questions proposed and (2) to build a machine learning model to predict flight prices.

1.1 Business questions

- Does price vary with Airlines?
- How is the price affected when tickets are bought in just 1 or 2 days before departure?
- Does ticket price change based on the departure time and arrival time?
- How the price changes with change in Source and Destination?
- How does the ticket price vary between Economy and Business class?

2 Data Understanding

2.1 Libraries import and loading the data

```
[ ]: # Data manipulation
import pandas as pd
import numpy as np

# EDA
import matplotlib.pyplot as plt
import seaborn as sns

# Machine Learning
from xgboost import XGBRegressor

# Pre-processing
import optuna
from sklearn import metrics
from sklearn.model_selection import train_test_split, KFold, cross_val_score
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import FunctionTransformer
from sklearn.compose import ColumnTransformer
```

```

from sklearn.impute import SimpleImputer
from category_encoders import TargetEncoder

# Notebook config
import warnings
warnings.filterwarnings('ignore')
plt.style.use('ggplot')
palette = ["#1F4E79", "#4F81BD", "#A9CCE3", "#D9EAD3", "#BFBFBF", "#595959",
           ↪ "#D0E0EB", "#8FA3BF"]
sns.set_palette(palette)
pd.set_option('display.max_rows', 50)
pd.set_option('display.max_columns', None)

```

```
[ ]: data = pd.read_csv("../data/raw/Clean_Dataset.csv")
```

2.2 Data check-up

```
[ ]: print(f"The dataset has {data.shape[0]} rows and {data.shape[1]} columns.")
```

The dataset has 300153 rows and 12 columns.

```
[ ]: print(f"The dataset has: {len(data.select_dtypes(include = 'object').columns.
           ↪to_list())} categorical columns.")
print(f"The dataset has: {len(data.select_dtypes(include = 'number').columns.
           ↪to_list())} numeric columns.")
```

The dataset has: 8 categorical columns.

The dataset has: 4 numeric columns.

```
[ ]: data.nunique().sort_values()
```

```
[ ]: class          2
      stops         3
      airline       6
      source_city   6
      departure_time 6
      arrival_time  6
      destination_city 6
      days_left     49
      duration      476
      flight        1561
      price         12157
      Unnamed: 0    300153
      dtype: int64
```

We've got to deleted this Unnamed: 0 column (probably the index)

```
[ ]: data.isnull().mean()
```

```
[ ]: Unnamed: 0      0.0
      airline        0.0
      flight         0.0
      source_city    0.0
      departure_time 0.0
      stops          0.0
      arrival_time   0.0
      destination_city 0.0
      class          0.0
      duration       0.0
      days_left      0.0
      price          0.0
      dtype: float64
```

We've got no null data.

```
[ ]: data.dtypes
```

```
[ ]: Unnamed: 0      int64
      airline        object
      flight         object
      source_city    object
      departure_time  object
      stops          object
      arrival_time   object
      destination_city object
      class          object
      duration       float64
      days_left      int64
      price          int64
      dtype: object
```

```
[ ]: data.head()
```

```
[ ]: Unnamed: 0  airline  flight  source_city  departure_time  stops  \
0            0  SpiceJet  SG-8709      Delhi      Evening    zero
1            1  SpiceJet  SG-8157      Delhi  Early_Morning    zero
2            2  AirAsia   I5-764      Delhi  Early_Morning    zero
3            3  Vistara   UK-995      Delhi      Morning    zero
4            4  Vistara   UK-963      Delhi      Morning    zero

      arrival_time  destination_city  class  duration  days_left  price
0           Night      Mumbai  Economy     2.17         1    5953
1          Morning      Mumbai  Economy     2.33         1    5953
2  Early_Morning      Mumbai  Economy     2.17         1    5956
3       Afternoon      Mumbai  Economy     2.25         1    5955
4          Morning      Mumbai  Economy     2.33         1    5955
```

We don't need further data manipulation besides excluding the columns `Unnamed: 0` and `flight`

```
[ ]: df = data.drop(columns = ['Unnamed: 0', 'flight'], axis = 1).copy()
df.head()
```

```
[ ]:      airline source_city departure_time stops  arrival_time destination_city \
0  SpiceJet      Delhi      Evening zero      Night      Mumbai
1  SpiceJet      Delhi  Early_Morning zero      Morning      Mumbai
2  AirAsia      Delhi  Early_Morning zero  Early_Morning      Mumbai
3  Vistara      Delhi      Morning zero      Afternoon      Mumbai
4  Vistara      Delhi      Morning zero      Morning      Mumbai

      class  duration  days_left  price
0  Economy      2.17          1  5953
1  Economy      2.33          1  5953
2  Economy      2.17          1  5956
3  Economy      2.25          1  5955
4  Economy      2.33          1  5955
```

2.3 EDA

```
[ ]: df.select_dtypes(include = 'object').describe().T
```

```
[ ]:      count unique    top    freq
airline    300153      6  Vistara  127859
source_city 300153      6   Delhi   61343
departure_time 300153      6  Morning   71146
stops       300153      3     one  250863
arrival_time 300153      6   Night   91538
destination_city 300153      6  Mumbai   59097
class       300153      2  Economy  206666
```

```
[ ]: df.describe().round(2).T
```

```
[ ]:      count      mean      std      min      25%      50%      75% \
duration  300153.0    12.22      7.19      0.83      6.83     11.25     16.17
days_left 300153.0    26.00     13.56      1.00     15.00     26.00     38.00
price     300153.0  20889.66  22697.77  1105.00  4783.00  7425.00  42521.00

      max
duration    49.83
days_left   49.00
price       123071.00
```

```
[ ]: print(f"Target Skewness and Kurtosis")
print("=" * 30)
print(f"Price Skewness: {df['price'].skew():.4f}")
print(f"Price Kurtosis: {df['price'].kurt():.4f}")
```

Target Skewness and Kurtosis


=====

Price Skewness: 1.0614

Price Kurtosis: -0.3963

```
[ ]: fig, ax = plt.subplots(figsize = (12, 6))
sns.kdeplot(df.price, fill = True)

ax.set_title("Price Distribution", fontsize = 15, pad = 5, loc = 'left')
ax.set_xlabel("Price (INR)", fontsize = 8)
ax.set_ylabel("Frequency", fontsize = 8)
plt.show()
```




Our target is right-skewed, which may be caused by the price difference between the flight classes.

2.3.1 Does price vary with Airlines?

```
[ ]: fig, ax = plt.subplots(figsize = (12, 6))


sns.barplot(x = df.airline, y = df.price, color = '#1F4E79', order = df.
    ↳groupby('airline')['price'].mean().sort_values(ascending=False).index)
ax.set_title("Price x Airline", fontsize = 15, pad = 5, loc = 'left')
ax.set_xlabel("Airline", fontsize = 8)
ax.set_ylabel("Price (INR)", fontsize = 8)
plt.show()
```



2.3.2 Does price vary with class?

```
[ ]: fig, ax = plt.subplots(figsize = (12, 6))


sns.barplot(x = df.airline, y = df.price, hue = df['class'], order = df.
    ↳groupby('airline')['price'].mean().sort_values(ascending=False).index)
ax.set_title("Price x Airline", fontsize = 15, pad = 5, loc = 'left')
ax.set_xlabel("Airline", fontsize = 8)
ax.set_ylabel("Price (INR)", fontsize = 8)
ax.legend(title = 'Class')
plt.show()
```



Vistara and Air India have higher prices, but does it remain this way when looking only at economy class flights?

```
[ ]: econ = df[df['class'] == 'Economy']

fig, ax = plt.subplots(figsize = (12, 6))
sns.barplot(x = econ.airline, y = econ.price, color = '#1F4E79', order = econ.
    ↳groupby('airline')['price'].mean().sort_values(ascending=False).index)
ax.set_title("Price x Airline in Economy Class", fontsize = 15, pad = 5, loc =_
    ↳'left')
ax.set_xlabel("Airline", fontsize = 8)
ax.set_ylabel("Price (INR)", fontsize = 8)
plt.show()
```



Vistara and Air India still have the higher prices.

2.3.3 How is the price affected when tickets are bought in just 1 or 2 days before departure?

```
[ ]: def days(i):
    if i > 2:
        return 'Over 2 days'
    else:
        return 'Under 2 days'
```

```
[ ]: df_days = df.copy()
df_days['days'] = df['days_left'].apply(days)
```

```
[ ]: days = df_days.groupby("days").agg(price_mean = ('price', 'mean'), price_median_
    => ('price', 'median')).reset_index()
days
```

```
[ ]:
    days    price_mean  price_median
0  Over 2 days  20757.498484      7347.0
1  Under 2 days  27421.169326     16739.0
```

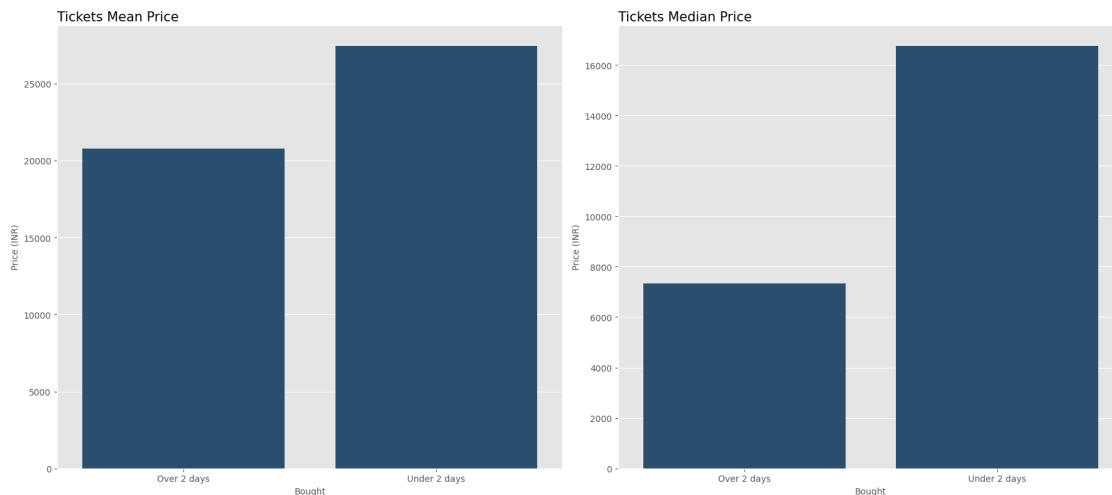
```
[ ]: fig, axes = plt.subplots(1, 2, figsize = (18, 8))

sns.barplot(ax = axes[0], x = days.days, y = days.price_mean)
axes[0].set_title("Tickets Mean Price", fontsize = 15, pad = 5, loc = 'left')
axes[0].set_xlabel("Bought", fontsize = 10)
axes[0].set_ylabel("Price (INR)", fontsize = 10)
```



```
sns.barplot(ax = axes[1], x = days.days, y = days.price_median)
axes[1].set_title("Tickets Median Price", fontsize = 15, pad = 5, loc = 'left')
axes[1].set_xlabel("Bought", fontsize = 10)
axes[1].set_ylabel("Price (INR)", fontsize = 10)

plt.tight_layout()
plt.show()
```



The price of airline tickets tends to be higher the closer you are to the flight date.

2.3.4 Does ticket price change based on the departure time and arrival time?

```
[ ]: departure = df.groupby(["departure_time"]).agg(price_mean = ('price', 'mean'),
    ↳ price_median = ('price', 'median')).reset_index()
departure
```

```
[ ]: departure_time  price_mean  price_median
0      Afternoon    18179.203331      6663.0
1  Early_Morning    20370.676718      7212.0
2       Evening    21232.361894      7425.0
3   Late_Night     9295.299387      4499.0
4       Morning    21630.760254      8112.0
5         Night    23062.146808      7813.0
```

```
[ ]: fig, axes = plt.subplots(1, 2, figsize = (18, 8))

sns.barplot(ax = axes[0], x = departure.departure_time, y = departure.
    ↳ price_mean, color = '#1F4E79', order = departure.departure_time)
axes[0].set_title("Tickets Mean Price", fontsize = 15, pad = 5, loc = 'left')
```


```

axes[0].set_xlabel("Departure Time", fontsize = 10)
axes[0].set_ylabel("Price (INR)", fontsize = 10)

sns.barplot(ax = axes[1], x = departure.departure_time, y = departure.
    ↳price_median, color = '#1F4E79', order = departure.departure_time)
axes[1].set_title("Tickets Median Price", fontsize = 15, pad = 5, loc = 'left')
axes[1].set_xlabel("Departure Time", fontsize = 10)
axes[1].set_ylabel("Price (INR)", fontsize = 10)

plt.tight_layout()
plt.show()

```



Morning and night have the highest prices in departure time.

```

[ ]: arrival = df.groupby(["arrival_time"]).agg(price_mean = ('price', 'mean'),
    ↳price_median = ('price', 'median')).reset_index()
arrival

```

```

[ ]:
   arrival_time  price_mean  price_median
0   Afternoon    18494.598993         6714.0
1  Early_Morning    14993.139521         5800.0
2    Evening     23044.371615         8854.0
3   Late_Night    11284.906078         4867.0
4    Morning     22231.076098         7687.0
5     Night     21586.758341         7584.0

```

```

[ ]: fig, axes = plt.subplots(1, 2, figsize = (18, 8))

sns.barplot(ax = axes[0], x = arrival.arrival_time, y = arrival.price_mean,
    ↳color = '#1F4E79', order = arrival.arrival_time)

```

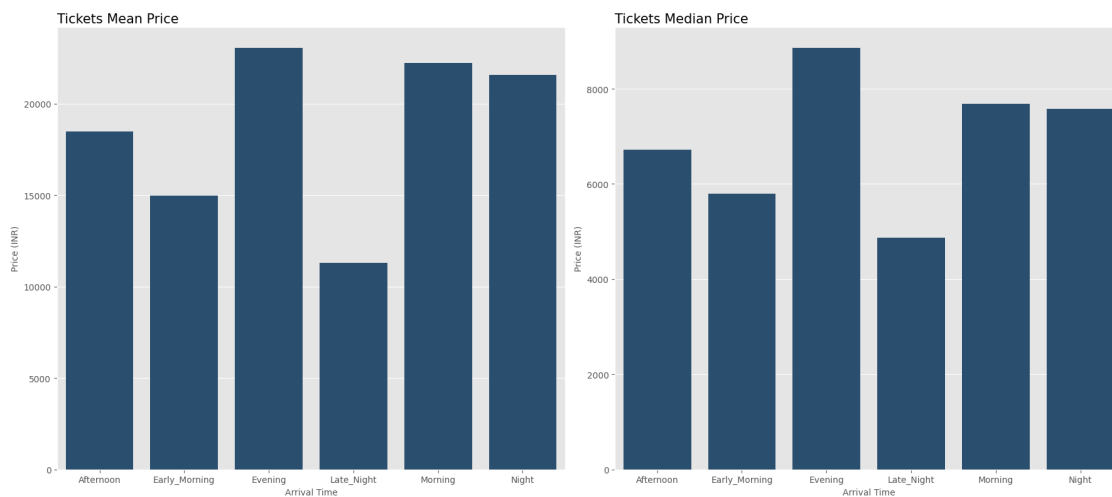
```

axes[0].set_title("Tickets Mean Price", fontsize = 15, pad = 5, loc = 'left')
axes[0].set_xlabel("Arrival Time", fontsize = 10)
axes[0].set_ylabel("Price (INR)", fontsize = 10)

sns.barplot(ax = axes[1], x = arrival.arrival_time, y = arrival.price_median,
            color = '#1F4E79', order = arrival.arrival_time)
axes[1].set_title("Tickets Median Price", fontsize = 15, pad = 5, loc = 'left')
axes[1].set_xlabel("Arrival Time", fontsize = 10)
axes[1].set_ylabel("Price (INR)", fontsize = 10)

plt.tight_layout()
plt.show()

```



Evening and morning have the highest prices in arrival time.

2.3.5 How the price changes with change in Source and Destination?

```

[ ]: source = df.groupby("source_city").agg(price_mean = ('price', 'mean'),
            price_median = ('price', 'median')).reset_index()
source

```

```

[ ]:
source_city  price_mean  price_median
0    Bangalore  21469.460575      7488.0
1     Chennai  21995.339871      7846.0
2      Delhi  18951.326639      6840.0
3  Hyderabad  20155.623879      6855.0
4    Kolkata  21746.235679      7958.0
5     Mumbai  21483.818839      7413.0

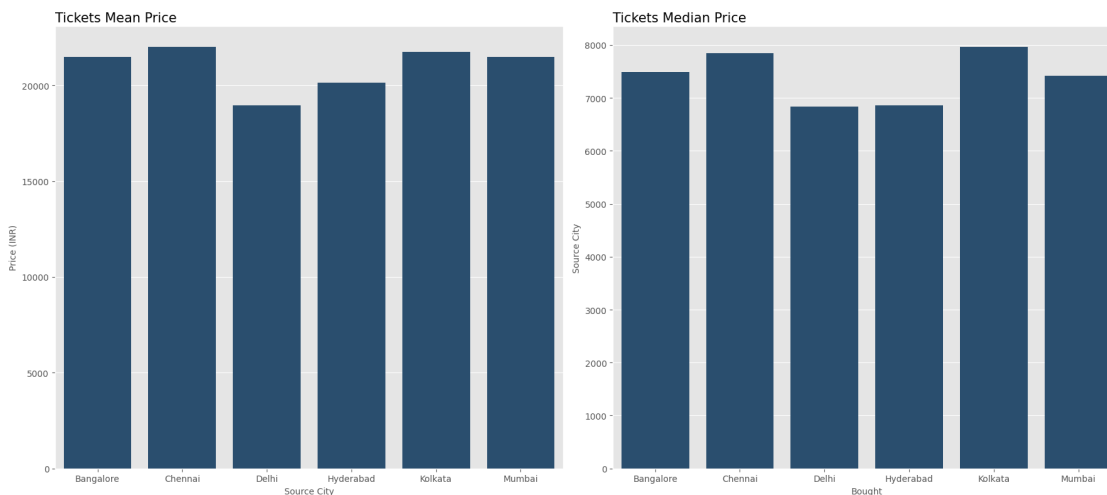
```

```
[ ]: fig, axes = plt.subplots(1, 2, figsize = (18, 8))

sns.barplot(ax = axes[0], x = source.source_city, y = source.price_mean, color='
↳ '#1F4E79', order = source.source_city)
axes[0].set_title("Tickets Mean Price", fontsize = 15, pad = 5, loc = 'left')
axes[0].set_xlabel("Source City", fontsize = 10)
axes[0].set_ylabel("Price (INR)", fontsize = 10)

sns.barplot(ax = axes[1], x = source.source_city, y = source.price_median,
↳ color = '#1F4E79', order = source.source_city)
axes[1].set_title("Tickets Median Price", fontsize = 15, pad = 5, loc = 'left')
axes[1].set_xlabel("Bought", fontsize = 10)
axes[1].set_ylabel("Source City", fontsize = 10)

plt.tight_layout()
plt.show()
```



```
[ ]: destination = df.groupby("destination_city").agg(price_mean = ('price',
↳ 'mean'), price_median = ('price', 'median')).reset_index()
destination
```

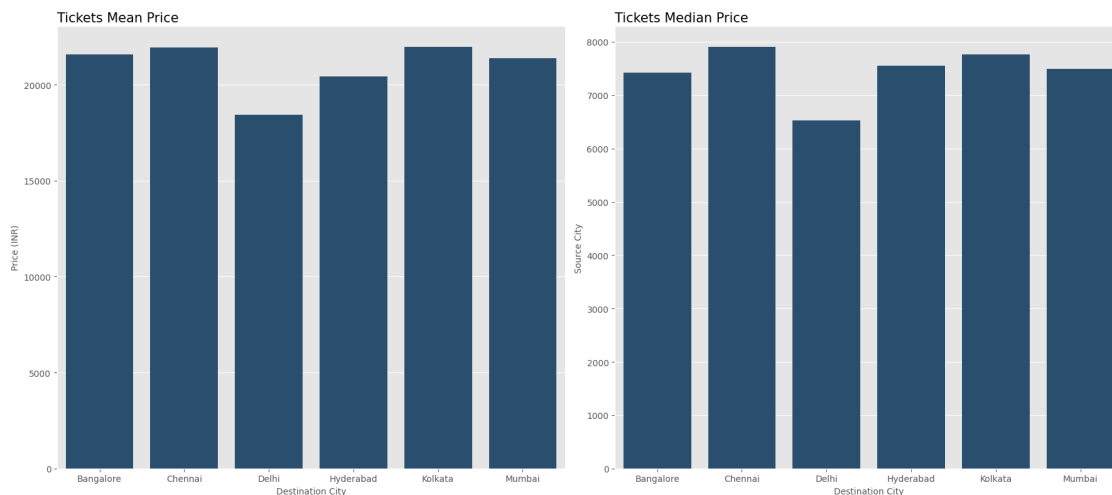
```
[ ]: destination_city  price_mean  price_median
0      Bangalore    21593.955784      7425.0
1      Chennai     21953.323969      7900.0
2      Delhi       18436.767870      6521.0
3      Hyderabad   20427.661284      7548.0
4      Kolkata     21959.557556      7767.0
5      Mumbai      21372.529469      7496.0
```

```
[ ]: fig, axes = plt.subplots(1, 2, figsize = (18, 8))

sns.barplot(ax = axes[0], x = destination.destination_city, y = destination.
    price_mean, color = '#1F4E79', order = destination.destination_city)
axes[0].set_title("Tickets Mean Price", fontsize = 15, pad = 5, loc = 'left')
axes[0].set_xlabel("Destination City", fontsize = 10)
axes[0].set_ylabel("Price (INR)", fontsize = 10)

sns.barplot(ax = axes[1], x = destination.destination_city, y = destination.
    price_median, color = '#1F4E79', order = destination.destination_city)
axes[1].set_title("Tickets Median Price", fontsize = 15, pad = 5, loc = 'left')
axes[1].set_xlabel("Destination City", fontsize = 10)
axes[1].set_ylabel("Source City", fontsize = 10)

plt.tight_layout()
plt.show()
```



Kolkata and Chennai have the most expensive tickets.

2.3.6 Correlation Matrix

```
[ ]: dummies = pd.get_dummies(df, prefix = ['airline',
    'source_city',
    'departure_time',
    'stops',
    'arrival_time',
    'destination_city',
    'class'], dtype = int)
corr = dummies.corr()
```

```
[ ]: corr.loc['price'].sort_values(ascending = False)
```

```
[ ]: price
      class_Business      0.937860
      airline_Vistara      0.360816
      duration             0.204222
      stops_one            0.199913
      airline_Air_India     0.070041
      arrival_time_Evening   0.056408
      departure_time_Night   0.041768
      arrival_time_Morning   0.030379
      destination_city_Kolkata 0.020956
      arrival_time_Night     0.020344
      source_city_Chennai    0.018742
      destination_city_Chennai 0.018473
      departure_time_Morning 0.018199
      source_city_Kolkata    0.016127
      destination_city_Bangalore 0.014050
      source_city_Mumbai     0.013206
      source_city_Bangalore  0.011702
      destination_city_Mumbai 0.010533
      departure_time_Evening 0.007946
      destination_city_Hyderabad -0.008292
      departure_time_Early_Morning -0.012232
      source_city_Hyderabad -0.012828
      departure_time_Late_Night -0.033768
      arrival_time_Afternoon -0.040258
      source_city_Delhi      -0.043282
      departure_time_Afternoon -0.051968
      destination_city_Delhi -0.052527
      arrival_time_Early_Morning -0.060449
      stops_two_or_more      -0.064248
      days_left              -0.091949
      arrival_time_Late_Night -0.093602
      airline_SpiceJet        -0.114019
      airline_AirAsia         -0.176188
      stops_zero              -0.187277
      airline_GO_FIRST        -0.194179
      airline_Indigo          -0.280882
      class_Economy           -0.937860
      Name: price, dtype: float64
```


The highlights:

- Business class has a quite high correlation with price.
- Vistara has the highest positive correlation with price, which is the highest among the airlines.
- Duration and single stop also have significant correlations.

Why not visualize this?

```
[ ]: fig = plt.figure(figsize = (6, 12))
ax = fig.add_axes([0, 0, 1, 1])
sns.heatmap(corr[['price']].sort_values(by = 'price', ascending = False), annot_
↵= True, fmt = '.4f', cmap = 'Blues', linecolor='black', linewidths=0.5)
ax.set_title("Correlation Matrix", fontsize = 18, pad = 10, loc = 'left')
plt.show()
```

Correlation Matrix



3 Data Preparation

```
[ ]: features = df.drop(columns = ['price'], axis = 1).columns.to_list()
    target = 'price'

    X = df[features]
    y = df[target]

[ ]: X_temp, X_test, y_temp, y_test = train_test_split(X, y, test_size=0.20,
    ↪random_state=21)

    X_train, X_val, y_train, y_val = train_test_split(X_temp, y_temp, test_size = 0.
    ↪25, random_state=21)

[ ]: cat_features = X_train.select_dtypes(exclude = 'number').columns.to_list()
    num_features = X_train.select_dtypes(include = 'number').columns.to_list()

[ ]: cat_transformer = Pipeline([
    ('imput', SimpleImputer(strategy = "most_frequent")),
    ('enconder', TargetEncoder())
])

    num_transformer = Pipeline([
    ('imput', SimpleImputer(strategy = 'median')),
    ])

    preprocessor = ColumnTransformer(
        transformers = [
            ('cat', cat_transformer, cat_features),
            ('num', num_transformer, num_features)
        ])
    ])
```

4 Modeling

4.1 The base model

```
[ ]: model = XGBRegressor(objective='reg:squarederror', random_state = 21)

    reg = Pipeline([
    ('preprocessor', preprocessor),
    ('regressor', model)
    ])

    reg.fit(X_train, y_train)
```

```
[ ]: Pipeline(steps=[('preprocessor',
                      ColumnTransformer(transformers=[('cat',
                                                       Pipeline(steps=[('imputer',
                                                                           SimpleImputer(strategy='most_frequent')),
                                                                           TargetEncoder()))]),
                      ['airline', 'source_city',
                       'departure_time', 'stops',
                       'arrival_time',
                       'destination_city',
                       'class']),
                  ('num',
                   Pipeline(steps=[('imputer',
                                     SimpleImputer(strategy='median'))]),
                   ['duration', 'days_le...
feature_types=None, gamma=None, grow_policy=None,
importance_type=None,
interaction_constraints=None, learning_rate=None,
max_bin=None, max_cat_threshold=None,
max_cat_to_onehot=None, max_delta_step=None,
max_depth=None, max_leaves=None,
min_child_weight=None, missing=nan,
monotone_constraints=None, multi_strategy=None,
n_estimators=None, n_jobs=None,
num_parallel_tree=None, random_state=21, ...))])
```

```
[ ]: y_pred = reg.predict(X_val)
```

```
[ ]: print(f"Validation Set")
print("=" * 30)
print(f"MSE: {metrics.mean_squared_error(y_val, y_pred):.4f}")
print(f"RMSE: {metrics.mean_squared_error(y_val, y_pred, squared = False):.4f}")
print(f"MAE: {metrics.mean_absolute_error(y_val, y_pred):.4f}")
print(f"MAPE: {metrics.mean_absolute_percentage_error(y_val, y_pred):.4f}")
print(f"R2 Score: {metrics.r2_score(y_val, y_pred):.4f}")
```

Validation Set

=====

MSE: 12790568.0237

RMSE: 3576.3904

MAE: 2048.7408

MAPE: 0.1504

R2 Score: 0.9751

These are the base model results let's work it better with hyperparamer optimization.

4.2 Optimized model

```
[ ]: def objective(trial):  
    model = Pipeline([  
        ('preprocessor', preprocessor),  
        ('regressor', XGBRegressor(  
            objective = 'reg:squarederror',  
            n_estimators = trial.suggest_int("n_estimators", 100, 1000),  
            learning_rate = trial.suggest_float("learning_rate", 1e-3, 0.1,   
↪ log=True),  
            max_depth = trial.suggest_int("max_depth", 1, 10),  
            subsample = trial.suggest_float("subsample", 0.05, 1.0),  
            colsample_bytree = trial.suggest_float("colsample_bytree", 0.05, 1.0),  
            min_child_weight = trial.suggest_int("min_child_weight", 1, 20),  
            random_state = 21  
        ))  
    ])  
  
    model.fit(X_train, y_train)  
    predictions = model.predict(X_val)  
    rmse = metrics.mean_squared_error(y_val, predictions, squared=False)  
    return rmse
```

```
[ ]: study = optuna.create_study(direction='minimize')  
study.optimize(objective, n_trials=30)
```

```
[I 2024-08-25 18:59:56,491] A new study created in memory with name: no-  
name-77cb1f25-546f-450e-b906-570e0a907ace  
[I 2024-08-25 19:00:02,407] Trial 0 finished with value: 8862.57658862416 and  
parameters: {'n_estimators': 229, 'learning_rate': 0.024492389628324486,  
'max_depth': 6, 'subsample': 0.2550588484785341, 'colsample_bytree':  
0.23487901404061928, 'min_child_weight': 9}. Best is trial 0 with value:  
8862.57658862416.  
[I 2024-08-25 19:00:12,848] Trial 1 finished with value: 13377.321509158486 and  
parameters: {'n_estimators': 833, 'learning_rate': 0.0011862800787346272,  
'max_depth': 2, 'subsample': 0.287177984250061, 'colsample_bytree':  
0.597102663039611, 'min_child_weight': 19}. Best is trial 0 with value:  
8862.57658862416.  
[I 2024-08-25 19:00:17,781] Trial 2 finished with value: 10143.12062061122 and  
parameters: {'n_estimators': 363, 'learning_rate': 0.005012523465267334,  
'max_depth': 1, 'subsample': 0.9066574973209431, 'colsample_bytree':  
0.5745271035798778, 'min_child_weight': 6}. Best is trial 0 with value:  
8862.57658862416.  
[I 2024-08-25 19:00:25,659] Trial 3 finished with value: 5220.244817019831 and  
parameters: {'n_estimators': 578, 'learning_rate': 0.013331414719213857,  
'max_depth': 3, 'subsample': 0.7338248812020122, 'colsample_bytree':  
0.5528055750913233, 'min_child_weight': 17}. Best is trial 3 with value:  
5220.244817019831.
```

[I 2024-08-25 19:00:40,638] Trial 4 finished with value: 3430.148895520396 and parameters: {'n_estimators': 835, 'learning_rate': 0.034093230314878084, 'max_depth': 7, 'subsample': 0.8425699635909832, 'colsample_bytree': 0.7639393898246126, 'min_child_weight': 1}. Best is trial 4 with value: 3430.148895520396.

[I 2024-08-25 19:00:46,678] Trial 5 finished with value: 4345.345818433851 and parameters: {'n_estimators': 362, 'learning_rate': 0.03420095978797079, 'max_depth': 5, 'subsample': 0.3027445461141159, 'colsample_bytree': 0.7972863298661383, 'min_child_weight': 2}. Best is trial 4 with value: 3430.148895520396.

[I 2024-08-25 19:01:06,309] Trial 6 finished with value: 4098.152181211634 and parameters: {'n_estimators': 723, 'learning_rate': 0.004478972337335129, 'max_depth': 10, 'subsample': 0.9377742222228224, 'colsample_bytree': 0.7835008738456918, 'min_child_weight': 15}. Best is trial 4 with value: 3430.148895520396.

[I 2024-08-25 19:01:16,206] Trial 7 finished with value: 3900.9973631500907 and parameters: {'n_estimators': 496, 'learning_rate': 0.018072673176409943, 'max_depth': 7, 'subsample': 0.19289290264290615, 'colsample_bytree': 0.8938799879780247, 'min_child_weight': 4}. Best is trial 4 with value: 3430.148895520396.

[I 2024-08-25 19:01:21,297] Trial 8 finished with value: 18818.848514541034 and parameters: {'n_estimators': 186, 'learning_rate': 0.003870787437535607, 'max_depth': 7, 'subsample': 0.9370731995203976, 'colsample_bytree': 0.3177843378411038, 'min_child_weight': 20}. Best is trial 4 with value: 3430.148895520396.

[I 2024-08-25 19:01:27,635] Trial 9 finished with value: 18099.112077715938 and parameters: {'n_estimators': 428, 'learning_rate': 0.001346066777203931, 'max_depth': 3, 'subsample': 0.16951403572568063, 'colsample_bytree': 0.33594353062755317, 'min_child_weight': 13}. Best is trial 4 with value: 3430.148895520396.

[I 2024-08-25 19:01:59,920] Trial 10 finished with value: 2643.9141673555478 and parameters: {'n_estimators': 1000, 'learning_rate': 0.09966133771268752, 'max_depth': 10, 'subsample': 0.5688373909283602, 'colsample_bytree': 0.9316787148859975, 'min_child_weight': 9}. Best is trial 10 with value: 2643.9141673555478.

[I 2024-08-25 19:02:26,792] Trial 11 finished with value: 2622.1551841132173 and parameters: {'n_estimators': 956, 'learning_rate': 0.08885057738863303, 'max_depth': 10, 'subsample': 0.560743995211729, 'colsample_bytree': 0.9985920323190209, 'min_child_weight': 9}. Best is trial 11 with value: 2622.1551841132173.

[I 2024-08-25 19:02:56,855] Trial 12 finished with value: 2641.5514477161096 and parameters: {'n_estimators': 982, 'learning_rate': 0.09640189438514946, 'max_depth': 10, 'subsample': 0.554626170108951, 'colsample_bytree': 0.9697152090965963, 'min_child_weight': 10}. Best is trial 11 with value: 2622.1551841132173.

[I 2024-08-25 19:03:19,096] Trial 13 finished with value: 2714.4756414577146 and parameters: {'n_estimators': 965, 'learning_rate': 0.082850949744856, 'max_depth': 9, 'subsample': 0.4668061915390079, 'colsample_bytree':

0.9930249947994891, 'min_child_weight': 12}. Best is trial 11 with value: 2622.1551841132173.

[I 2024-08-25 19:03:27,477] Trial 14 finished with value: 6632.783312487924 and parameters: {'n_estimators': 670, 'learning_rate': 0.05417508190052682, 'max_depth': 9, 'subsample': 0.5511198035324097, 'colsample_bytree': 0.10550187167886843, 'min_child_weight': 7}. Best is trial 11 with value: 2622.1551841132173.

[I 2024-08-25 19:03:45,165] Trial 15 finished with value: 2790.9313417747553 and parameters: {'n_estimators': 870, 'learning_rate': 0.05247343772861899, 'max_depth': 9, 'subsample': 0.6958264863288773, 'colsample_bytree': 0.6780883799235097, 'min_child_weight': 11}. Best is trial 11 with value: 2622.1551841132173.

[I 2024-08-25 19:04:00,367] Trial 16 finished with value: 3852.8169513821376 and parameters: {'n_estimators': 720, 'learning_rate': 0.008591745322137117, 'max_depth': 8, 'subsample': 0.4434147021134358, 'colsample_bytree': 0.8924115303253563, 'min_child_weight': 14}. Best is trial 11 with value: 2622.1551841132173.

[I 2024-08-25 19:04:13,996] Trial 17 finished with value: 3735.6769928667954 and parameters: {'n_estimators': 903, 'learning_rate': 0.06092789943867181, 'max_depth': 5, 'subsample': 0.6527694236750965, 'colsample_bytree': 0.9801787225334296, 'min_child_weight': 8}. Best is trial 11 with value: 2622.1551841132173.

[I 2024-08-25 19:04:36,304] Trial 18 finished with value: 3184.1916607492976 and parameters: {'n_estimators': 777, 'learning_rate': 0.03225538748462477, 'max_depth': 10, 'subsample': 0.07067098169689312, 'colsample_bytree': 0.6752219938856643, 'min_child_weight': 5}. Best is trial 11 with value: 2622.1551841132173.

[I 2024-08-25 19:04:52,092] Trial 19 finished with value: 4767.629480219714 and parameters: {'n_estimators': 609, 'learning_rate': 0.009562616898461822, 'max_depth': 8, 'subsample': 0.38756963664584776, 'colsample_bytree': 0.45425041049110376, 'min_child_weight': 10}. Best is trial 11 with value: 2622.1551841132173.

[I 2024-08-25 19:05:14,242] Trial 20 finished with value: 2717.2916679968557 and parameters: {'n_estimators': 978, 'learning_rate': 0.09451574613160346, 'max_depth': 8, 'subsample': 0.8083091583103925, 'colsample_bytree': 0.833392984731383, 'min_child_weight': 16}. Best is trial 11 with value: 2622.1551841132173.

[I 2024-08-25 19:05:53,105] Trial 21 finished with value: 2634.378721971906 and parameters: {'n_estimators': 990, 'learning_rate': 0.09583541671831836, 'max_depth': 10, 'subsample': 0.5969909240808767, 'colsample_bytree': 0.9184009108796035, 'min_child_weight': 10}. Best is trial 11 with value: 2622.1551841132173.

[I 2024-08-25 19:06:25,054] Trial 22 finished with value: 2635.969874375013 and parameters: {'n_estimators': 909, 'learning_rate': 0.06044884059919022, 'max_depth': 10, 'subsample': 0.6212571409729032, 'colsample_bytree': 0.9958393871003632, 'min_child_weight': 11}. Best is trial 11 with value: 2622.1551841132173.

[I 2024-08-25 19:06:50,515] Trial 23 finished with value: 2742.4568220592378 and

```

parameters: {'n_estimators': 864, 'learning_rate': 0.05955187642511468,
'max_depth': 9, 'subsample': 0.6213721417231871, 'colsample_bytree':
0.8606082660923033, 'min_child_weight': 12}. Best is trial 11 with value:
2622.1551841132173.
[I 2024-08-25 19:07:19,748] Trial 24 finished with value: 2663.65379376364 and
parameters: {'n_estimators': 915, 'learning_rate': 0.043579686126519926,
'max_depth': 10, 'subsample': 0.7550692310656586, 'colsample_bytree':
0.7210719048107044, 'min_child_weight': 7}. Best is trial 11 with value:
2622.1551841132173.
[I 2024-08-25 19:07:38,849] Trial 25 finished with value: 3084.2276678095877 and
parameters: {'n_estimators': 778, 'learning_rate': 0.022136167906717383,
'max_depth': 9, 'subsample': 0.46755345479520605, 'colsample_bytree':
0.9049239772672324, 'min_child_weight': 11}. Best is trial 11 with value:
2622.1551841132173.
[I 2024-08-25 19:07:56,193] Trial 26 finished with value: 2863.4264035442498 and
parameters: {'n_estimators': 783, 'learning_rate': 0.06757966034251568,
'max_depth': 8, 'subsample': 0.622856073631866, 'colsample_bytree':
0.9952279668630203, 'min_child_weight': 13}. Best is trial 11 with value:
2622.1551841132173.
[I 2024-08-25 19:08:12,012] Trial 27 finished with value: 3545.18621684338 and
parameters: {'n_estimators': 916, 'learning_rate': 0.044795768275337454,
'max_depth': 6, 'subsample': 0.39328505913497, 'colsample_bytree':
0.8193048186197822, 'min_child_weight': 8}. Best is trial 11 with value:
2622.1551841132173.
[I 2024-08-25 19:08:38,832] Trial 28 finished with value: 5420.144191965624 and
parameters: {'n_estimators': 925, 'learning_rate': 0.002147573427269228,
'max_depth': 10, 'subsample': 0.6792164479192159, 'colsample_bytree':
0.9340497790264902, 'min_child_weight': 4}. Best is trial 11 with value:
2622.1551841132173.
[I 2024-08-25 19:08:48,815] Trial 29 finished with value: 4650.7264288550905 and
parameters: {'n_estimators': 662, 'learning_rate': 0.02858470121364258,
'max_depth': 4, 'subsample': 0.5024618000666115, 'colsample_bytree':
0.4560553241436307, 'min_child_weight': 9}. Best is trial 11 with value:
2622.1551841132173.

```

```

[ ]: print('Best hyperparameters:', study.best_params)
      print('Best RMSE:', study.best_value)

```

```

Best hyperparameters: {'n_estimators': 956, 'learning_rate':
0.08885057738863303, 'max_depth': 10, 'subsample': 0.560743995211729,
'colsample_bytree': 0.9985920323190209, 'min_child_weight': 9}
Best RMSE: 2622.1551841132173

```

```

[ ]: reg_optuna = Pipeline([
      ('preprocessor', preprocessor),
      ('regressor', XGBRegressor(**study.best_params, random_state = 21))
    ])

```

```
reg_optuna.fit(X_train, y_train)
```

```
[ ]: Pipeline(steps=[('preprocessor',
                      ColumnTransformer(transformers=[('cat',
                                                       Pipeline(steps=[('imput',
                                                                           SimpleImputer(strategy='most_frequent')),
                                                                           TargetEncoder()))],
                                           ['airline', 'source_city',
                                            'departure_time', 'stops',
                                            'arrival_time',
                                            'destination_city',
                                            'class']),
                      ('num',
                       Pipeline(steps=[('imput',
                                         SimpleImputer(strategy='median'))]),
                       ['duration', 'days_le...
                      feature_types=None, gamma=None, grow_policy=None,
                      importance_type=None,
                      interaction_constraints=None,
                      learning_rate=0.08885057738863303, max_bin=None,
                      max_cat_threshold=None, max_cat_to_onehot=None,
                      max_delta_step=None, max_depth=10,
                      max_leaves=None, min_child_weight=9, missing=nan,
                      monotone_constraints=None, multi_strategy=None,
                      n_estimators=956, n_jobs=None,
                      num_parallel_tree=None, random_state=21, ...))])
```

```
[ ]: y_pred_optuna = reg_optuna.predict(X_val)
```

```
[ ]: print("Validation Set - Optimized")
print("=" * 30)
print(f"MSE: {metrics.mean_squared_error(y_val, y_pred_optuna):.4f}")
print(f"RMSE: {metrics.mean_squared_error(y_val, y_pred_optuna, squared =False):.4f}")
print(f"MAE: {metrics.mean_absolute_error(y_val, y_pred_optuna):.4f}")
print(f"MAPE: {metrics.mean_absolute_percentage_error(y_val, y_pred_optuna):.4f}")
print(f"R2 Score: {metrics.r2_score(y_val, y_pred_optuna):.4f}")
```

Validation Set - Optimized

=====

MSE: 6875697.8096

RMSE: 2622.1552

MAE: 1333.2493

MAPE: 0.0971

R2 Score: 0.9866

Our final model got good metrics in the validation set. The Mean Absolute Percentage Error is around 9% which shows a good performance, time to evaluate using the test set.

5 Evaluation

5.1 The metrics

```
[ ]: y_pred_test = reg_optuna.predict(X_test)

[ ]: print("Test Set")
      print("=" * 30)
      print(f"MSE: {metrics.mean_squared_error(y_test, y_pred_test):.4f}")
      print(f"RMSE: {metrics.mean_squared_error(y_test, y_pred_test, squared = False):
            ↪.4f}")
      print(f"MAE: {metrics.mean_absolute_error(y_test, y_pred_test):.4f}")
      print(f"MAPE: {metrics.mean_absolute_percentage_error(y_test, y_pred_test):.
            ↪.4f}")
      print(f"R2 Score: {metrics.r2_score(y_test, y_pred_test):.4f}")
```

Test Set

=====

MSE: 6855205.6699

RMSE: 2618.2448

MAE: 1328.5562

MAPE: 0.0962

R2 Score: 0.9868

The model showed good metrics, a high R2 score, and good RMSE and MAE.

5.2 Cross-validation

```
[ ]: scoring = metrics.make_scorer(metrics.mean_absolute_error)
      cv = KFold(n_splits = 5, shuffle=True, random_state=21)

      scores = cross_val_score(reg_optuna, X_train, y_train, cv = cv, scoring =_
            ↪scoring)
      print(f'Mean MAE: {scores}')
      print(f'General Mean MAE: {scores.mean()}')
      print(f'MAE Standard Deviation: {scores.std()}')


Mean MAE: [1393.91755665 1378.22671365 1368.95018009 1383.35412869 1363.1419011
]
General Mean MAE: 1377.5180960373127
MAE Standard Deviation: 10.803590968524654
```

We can also see that the model demonstrated good generalization, with no signs of overfitting.

5.3 Visualization


```
[ ]: fig, ax = plt.subplots(figsize = (12, 6))

sns.scatterplot(x = y_test, y = y_pred_test)
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color = 'red')
ax.set_title(f"Real x Predito (R2 Score: {metrics.r2_score(y_test, y_pred_test):
↪.4f})", fontsize = 15, pad = 10, loc = 'left')
ax.set_xlabel("Real", fontsize = 8)
ax.set_ylabel("Previsto", fontsize = 8)
plt.show()
```



```
[ ]: residuals = y_test - y_pred_test

fig, ax = plt.subplots(figsize = (12, 6))
sns.kdeplot(residuals, fill = True)
ax.set_title("Residuals Distribution", fontsize = 15, pad = 5, loc = 'left')
ax.set_xlabel("Price (INR)", fontsize = 8)
ax.set_ylabel("Density", fontsize = 8)
plt.show()
```



6 Conclusions

That's my first time trying to work with the CRISP-DM framework, and it was a good experience throughout this project. Regarding the business questions, we can say that:

- Price varies with class and airline, especially with class. In the case of airlines, there is a significant difference when we consider Vistara and Air India;
- Kolkata and Chennai are the most expensive destinations;
- For arrival time, evening and morning have the highest prices, and for departure, morning and night do;
- Buying tickets in advance is the best way to get good prices.

For the model we got good generalization and results, for more of my work you can find me on: - [LinkedIn](#) - [Github](#)